


Theoretical Guarantees in Reinforcement Learning : Regret analysis

Shipra Agrawal

IEOR 8100 (Reinforcement Learning)



Quantifying the performance of RL algorithms

- Computational complexity
 - the amount of per-time-step computation the algorithm uses during learning;
- Space complexity
 - the amount of memory used by the algorithm;
- Learning complexity
 - a measure of how much experience the algorithm needs to learn in a given task.

Learning complexity analysis

- Using a single thread of experience
 - No resets or generative model

Vs.

Generative models (simulator)

- <http://hunch.net/~jl/projects/RL/EE.html>

Learning complexity analysis

- Optimal learning complexity?
 - “optimally explore” meaning to obtain maximum expected discounted reward $E[\sum_{t=1}^{\infty} \gamma^{t-1} r_t]$ over *a known prior* over the unknown MDPs
 - Tractable in special cases ($\gamma < 1$, *MAB*, Gittins, 1989).

PAC analysis: single thread of experience

- Bound number of steps on which suboptimal policy is played
- near-optimally on all but a polynomial number of steps (Kakade, 2003; Strehl and Littman, 2008b, Strehl et al 2006, 2009)
- Example: [Strehl et al 2006] (may been improved in recent work)

Theorem 1 *Let M be any MDP and let ϵ and δ be two positive real numbers. If Delayed Q-learning is executed on MDP M , then it will follow an ϵ -optimal policy on all but $O\left(\frac{SA}{(1-\gamma)^8 \epsilon^4} \ln \frac{1}{\delta} \ln \frac{1}{\epsilon(1-\gamma)} \ln \frac{SA}{\delta \epsilon(1-\gamma)}\right)$ timesteps, with probability at least $1 - \delta$.*

Generative model: Sample complexity for ϵ optimal

Some examples:

- An $O(SA)$ analysis of Q-learning.
 - Michael Kearns and Satinder Singh, Finite-Sample Convergence Rates for Q-learning and Indirect Algorithms NIPS 1999.
- An analysis of TD-lambda.
 - Michael Kearns and Satinder Singh, Bias-Variance Error Bounds for Temporal Difference Updates, COLT 2000.

Regret analysis


- Bound difference in total reward of algorithm compared to a benchmark policy
 - $Reg(M, T) = T \rho^*(s_0) - E[\sum_{t=1}^T r_t]$
 - $\rho^*(s_0)$ is infinite horizon average reward achieved by the best stationary policy π^*
- Episodic/single thread
 - Episodic: starting state is reset every H steps
 - Single thread: assume Communicating with diameter D
- Worst-case regret bounds; example
 - Theorem [Agrawal, Jia, 2017]: For **any communicating MDP** M with (unknown) diameter D, and for $T \geq S^5 A$, our algorithm achieves: $Regret(M, T) \leq \tilde{O}(D\sqrt{SAT})$
- Bayesian regret bounds: (expectation over known prior on MDP M); example
 - Theorem [Osband, Russo, Van Roy, 2013]: For **any known prior distribution** f on MDP M, our algorithm achieves in episodic setting: $E_{M \sim f}[Regret(M, T)] \leq \tilde{O}(HS\sqrt{AT})$

Lecture overview

- Exploration-exploitation and regret minimization for multi-armed bandit
 - UCB
 - Thompson Sampling
- Regret minimization for RL (tabular)
 - UCRL
 - Posterior sampling based algorithms

Multi-armed bandit == single state MDP

Two armed bandit



Stochastic multi-armed bandit problem

- Online decisions
 - At every time step $t = 1, \dots, T$, pull one arm out of N arms
- Stochastic feedback
 - For each arm i , reward is generated i.i.d. from a **fixed but unknown distribution** ν_i support $[0,1]$, mean μ_i
- Bandit feedback
 - Only the reward of the pulled arm can be observed



The multi-armed bandit problem (Thompson 1933; Robbins 1952)

Multiple rigged slot machines in a casino.

Which one to put money on?

- Try each one out

Arm == actions



WHEN TO STOP TRYING (EXPLORATION) AND START PLAYING (EXPLOITATION)?



Regret minimization problem

- Minimize expected regret in time T
 - $Regret(T) = T\mu^* - E[\sum_{t=1}^T r_t]$ where $\mu^* = \max_j \mu_j$
- Equivalent formulation
 - Expected regret for playing arm $I_t = i$ at time t : $\Delta_i = \mu^* - \mu_i$
 - $n_{i,T}$ be the number times arm i is played till time T
 - $Regret(T) = \sum_t (\mu^* - \mu_{I_t}) = \sum_{i:\mu_i \neq \mu^*} \Delta_i E[n_{i,T}]$
- If we can bound $n_{i,T}$ by $\frac{C \log(T)}{\Delta_i^2}$, then regret bound: $\sum_{i:\mu_i \neq \mu^*} \frac{C \log(T)}{\Delta_i}$
 - Problem-dependent bound, close to lower bound [Lai and Robbins 1985]
- For problem independent bound: $C\sqrt{NT \log(T)}$
 - Separately bound total regret for playing an arm with $\Delta_i \leq \sqrt{\frac{N \log(T)}{T}}$ and $\Delta_i > \sqrt{\frac{N \log(T)}{T}}$
 - Lower bound $\Omega(\sqrt{NT})$,



The need for exploration

- Two arms **black** and **red**
 - Random rewards with unknown mean $\mu_1 = 1.1$, $\mu_2 = 1$
 - Optimal expected reward in T time steps is $1.1 \times T$
- Exploit only strategy: use the current best estimate (MLE/empirical mean) of unknown mean to pick arms
- Initial few trials can mislead into playing red action forever

1.1, 1, 0.2,

1, 1, 1, 1, 1, 1, 1,

- Expected regret in T steps is close to $0.1 \times T$

Exploration-Exploitation tradeoff

- Exploitation: play the empirical mean reward maximizer
- Exploration: play less explored actions to ensure empirical estimates converge

Upper confidence bound algorithms for MAB

Optimism under face of uncertainty

- Empirical mean at time t for arm i

$$\hat{\mu}_{i,t} = \frac{\sum_{s=1: I_s=i}^t r_s}{n_{i,t}}$$

- Upper confidence bound

$$\text{UCB}_{i,t} := \hat{\mu}_{i,t} + 2\sqrt{\frac{\ln t}{n_{i,t}}}$$

Algorithm 1: UCB algorithm for the stochastic N-armed bandit problem

```
foreach  $t = 1, \dots, N$  do
  Play arm  $t$ 
end
foreach  $t = N + 1, N + 2, \dots, T$  do
  Play arm  $I_t = \arg \max_{i \in \{1 \dots N\}} \text{UCB}_{i,t-1}$ .
  Observe  $r_t$ , compute  $\text{UCB}_{i,t}$ 
end
```

Regret analysis

- Using Azuma-Hoeffding (since $E[r_t | I_t = i] = \mu_i$), with probability $1 - \frac{2}{t^2}$,

$$|\hat{\mu}_{i,t} - \mu_i| < \sqrt{\frac{4 \ln t}{n_{i,t}}}$$

- Two implications: for every arm i

- $UCB_{i,t} > \mu_i$ with probability $1 - \frac{2}{t^2}$

- If $n_{i,t} > \frac{16 \ln(T)}{\Delta_i^2}$, $\hat{\mu}_{i,t} < \mu_i + \frac{\Delta_i}{2}$

Recall: $UCB_{i,t} := \hat{\mu}_{i,t} + 2\sqrt{\frac{\ln t}{n_{i,t}}}$

- Each suboptimal arm is played at most $\frac{16 \ln(T)}{\Delta_i^2}$ times, since after that many plays:

$$\begin{aligned} UCB_{i,t} &= \hat{\mu}_{i,t} + \sqrt{\frac{\ln t}{n_{i,t}}} \leq \hat{\mu}_{i,t} + \frac{\Delta_i}{2} \\ &< \left(\mu_i + \frac{\Delta_i}{2} \right) + \frac{\Delta_i}{2} \\ &= \mu^* \\ &< UCB_{i^*,t} \end{aligned}$$


Thompson Sampling [Thompson, 1933]

- Natural and Efficient heuristic
- Maintain belief about parameters (e.g., mean reward) of each arm
- Observe feedback, update belief of pulled arm i in Bayesian manner
- Pull arm with posterior probability of being best arm
 - NOT same as choosing the arm that is most likely to be best



Posterior Sampling: main idea [Thompson 1933]

- Maintain Bayesian posteriors for unknown parameters
- With more trials posteriors concentrate on the true parameters
 - Mode captures MLE: enables exploitation
- Less trials means more uncertainty in estimates
 - Spread/variance captures uncertainty: enables exploration
- A sample from the posterior is used as an estimate for unknown parameters to make decisions




Bernoulli rewards, Beta priors

Uniform distribution $Beta(1,1)$

$Beta(\alpha, \beta)$ prior \Rightarrow Posterior

- $Beta(\alpha + 1, \beta)$ if you observe 1
- $Beta(\alpha, \beta + 1)$ if you observe 0



Start with $Beta(\alpha_0, \beta_0)$ prior belief for every arm

In round t ,

- For every arm i , sample $\theta_{i,t}$ independently from posterior $Beta(\alpha_0 + S_{i,t}, \beta_0 + F_{i,t})$
- Play arm $i_t = \max_i \theta_{i,t}$
- Observe reward and update the Beta posterior for arm i_t



Bayesian regret bounds for Thompson Sampling

- Theorem [Russo and Van Roy 2014]: Given any prior f over MAB instance M described by reward distributions with $[0,1]$ bounded support (e.g., prior distribution over $\mu_1, \mu_2, \dots, \mu_N$ in case of Bernoulli rewards)

$$\text{BayesianRegret}(T) = E_{M \sim f}[\text{Regret}(T, M)] \leq O(\sqrt{NT \ln T})$$

Note;

- $\text{Regret}(T, M)$ notation is used to explicitly indicate dependence on regret on the MAB instance M .
- In comparison, worst case regret,

$$\text{Regret}(T) = \max_M [\text{Regret}(T, M)]$$

Worst-case regret bounds

[A. and Goyal COLT 2012, AISTATS 2013, JACM 2017, Kaufmann et al. 2013]

Instance-dependent bounds for $\{0,1\}$ rewards

- $\text{Regret}(T) \leq \ln(T)(1 + \epsilon) \sum_i \frac{\Delta_i}{KL(\mu^* || \mu_i)} + O\left(\frac{N}{\epsilon^2}\right)$
 - Matches *asymptotic instance wise lower bound* [Lai Robbins 1985]
 - UCB algorithm achieves this only after careful tuning [Kaufmann et al. 2012]

Arbitrary bounded $[0,1]$ rewards (using Beta and Gaussian priors)

- $\text{Regret}(T) \leq O(\ln(T) \sum_i \frac{1}{\Delta_i})$
 - Matches the best available for UCB for general reward distributions

Instance-independent bounds (Beta and Gaussian priors)

- $\text{Regret}(T) \leq O(\sqrt{NT \ln T})$
- Prior and likelihood mismatch allowed!



Why does it work? Two arms example

- Two arms, $\mu_1 \geq \mu_2$, $\Delta = \mu_1 - \mu_2$
- Every time arm 2 is pulled, Δ regret
- ➡ ▪ Bound the number of pulls of arm 2 by $\frac{\log(T)}{\Delta^2}$ to get $\frac{\log(T)}{\Delta}$ regret bound
- How many pulls of arm 2 are actually needed?



Easy situation

After $n = O\left(\frac{\log(T)}{\Delta^2}\right)$ pulls of arm 2 **and arm 1**

- Empirical means are well separated

$$\text{Error } |\widehat{\mu}_i - \mu_i| \leq \sqrt{\frac{\log(T)}{n}} \leq \frac{\Delta}{4} \text{ whp}$$

(Using Azuma Hoeffding inequality)


- Beta Posteriors are well separated

$$\text{Mean} = \frac{\alpha_i}{\alpha_i + \beta_i} = \widehat{\mu}_i$$

$$\text{standard deviation} \simeq \frac{1}{\sqrt{\alpha + \beta}} = \frac{1}{\sqrt{n}} \leq \frac{\Delta}{4}$$

The two arms can be distinguished!

No more arm 2 pulls.




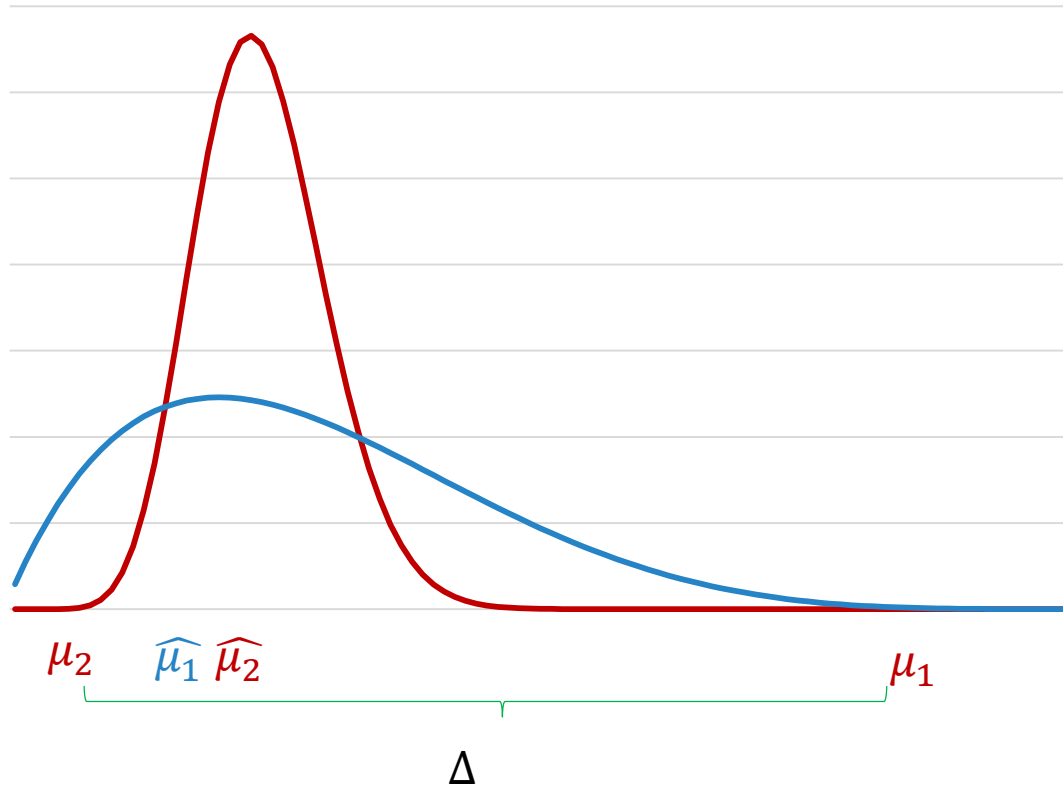
Easy situation

- If arm 2 is pulled less than $n = O\left(\frac{\log(T)}{\Delta^2}\right)$ times?
 - Regret is at most $n\Delta = \frac{\log(T)}{\Delta}$



Difficult situation

- At least $\frac{\log(T)}{\Delta^2}$ pulls of arm 2, but few pulls of arm 1



Main insight

- Arm 1 will be played roughly every constant number of steps in this situation
- It will take at most $constant \times \frac{\log T}{\Delta^2}$ steps (extra pulls of arm 2) to get out of this situation
- Total number of pulls of arm 2 is at most $O(\frac{\log T}{\Delta^2})$
- Summary: variance of posterior enables exploration
- Optimal bounds (and for multiple arms) require more careful use of posterior structure



Next...

Using UCB and TS ideas for exploration-exploitation in RL